Latent Trajectory Prediction for Grey Swan Events in Scarce-Label Multivariate Time Series

Anonymous Author(s)

Abstract

Multivariate time series from industrial, infrastructure, and scientific systems contain early warning signals for *grey swan events* - rare but partially predictable failures, anomalies, and threshold breaches that carry high stakes and few labels. We present a self-supervised method that predicts future sensor trajectories in latent space via a Joint Embedding Predictive Architecture (JEPA), yielding representations that support diverse grey swan tasks through trivial linear probes, and whose pretraining prediction error is itself a zero-label anomaly detector. On C-MAPSS FD001, a matched-architecture from-scratch ablation isolates the SSL contribution at $+8.8$ RMSE at full labels and $+16.9$ at 10%; at 5% labels the frozen encoder (21.53 ± 2.0) surpasses matched-budget supervised STAR (24.55 ± 6.4). An architectural ablation shows that *causal* context is the key inductive bias: a bidirectional encoder with VICReg regularisation transfers uniformly worse across FD002-FD004 ($+6$ to $+11$ RMSE) and collapses at low labels. On SMAP spacecraft telemetry, the prediction-error anomaly score attains PA-F1 = 62.5% with no anomaly labels (MTS-JEPA: 33.6%); non-PA F1 sits near the $\approx 7\%$ random baseline, so the detector is segment-localised rather than point-precise.

1 Introduction

A turbine blade cracks after 12,000 flight hours. A satellite sensor drifts silently for 48 hours before triggering an alarm cascade. The VIX doubles in a single trading session. An ICU patient develops sepsis six hours before clinical diagnosis. These are *grey swans*: events that are rare in operational data but follow partially predictable dynamics [Taleb, 2007, Lynch et al., 2025]. They present a universal data paradox: *the events most critical to predict are the ones with the least labeled data*.

Current approaches treat each event type as a separate problem with bespoke architectures: RUL prediction models [Fan et al., 2024, Li et al., 2025] never encounter anomaly detection benchmarks; anomaly detectors [Xu et al., 2022, Yang et al., 2023] never predict remaining useful life. Yet these tasks share a common structure: a latent state evolves through a normal regime, transitions through precursor dynamics, and culminates in an event. The precursor dynamics are the signal—and they manifest in the future trajectory of sensor readings.

We take this insight literally. Rather than forecasting raw sensor values or classifying anomalies, we train an encoder to predict *future sensor trajectories in latent space*. Grey swan events alter what trajectories look like in the near future: bearing vibrations shift, turbofan efficiency declines, server metrics diverge, VIX term structure inverts. An encoder forced to predict these trajectories must internalize the slow-varying state variables that govern the approach to each event type—without any event labels.

Trajectory JEPA. Our method is a causal Transformer encoder that predicts the latent representation of future observations $\mathbf{x}_{t+1:t+k}$ from past context $\mathbf{x}_{1:t}$, using an EMA target encoder [Grill et al., 2020, Assran et al., 2023]. Prediction in latent space avoids the reconstruction bias of masked autoencoders [He et al., 2022]—the encoder learns degradation dynamics, not signal-level detail.

Preprint.

One encoder, any event type. After pretraining, the frozen encoder supports *any* grey swan task through a trivial linear probe (Figure 2). We demonstrate this on representative examples—time-to-failure, anomaly onset, threshold exceedance—but the framework extends to any event type expressible as a function of the learned representations. A notable special case: the prediction error from pretraining is *itself* an anomaly detector, requiring zero event labels.

Contributions.

1. **Trajectory JEPA for grey swan prediction.** We adapt JEPA to multivariate time series via causal latent trajectory prediction, and demonstrate that a single pretrained encoder supports both time-to-failure regression (C-MAPSS) and zero-label anomaly onset detection (SMAP/MSL) through task-specific linear probes that share the same pretraining objective (sections 3, 3.4 and 5.2).
2. **Causal inductive bias matters more than model complexity.** A bidirectional context encoder with VICReg regularisation (V16b) attains comparable supervised fine-tuning RMSE but substantially worse frozen-probe generalisation (test RMSE 25.7 vs. 17.8), worse cross-machine transfer ($+6$ to $+11$ RMSE on FD002-FD004), and worse label efficiency at every budget (figs. 5 and 6, section 5.6).
3. **Quantified pretraining contribution and label-efficiency crossover.** A matched-architecture from-scratch ablation isolates the SSL objective at $+8.8$ RMSE at full labels, growing to $+16.9$ at 10% (section 5.3). At 5% labels, the frozen SSL encoder (21.53 ± 2.0) surpasses supervised STAR (24.55 ± 6.4), demonstrating the practical payoff in the scarce-label regime (section 5.3).
4. **Zero-label anomaly detection as a byproduct of pretraining.** The pretraining prediction error itself serves as an anomaly score without any event labels, achieving PA-F1 62.5% on SMAP (vs. 33.6% for MTS-JEPA) and 43.3% on MSL (section 3.4).

2 Related Work

Self-supervised learning for time series. Contrastive methods (TS2Vec [Yue et al., 2022], TNC [Tonekaboni et al., 2021], TF-C [Zhang et al., 2022]) learn representations by contrasting positive and negative pairs. Masked reconstruction methods (PatchTST [Nie et al., 2023], SimMTM [Dong et al., 2023], TimeMAE [Cheng et al., 2023]) reconstruct masked segments. JEPA avoids both: it predicts representations of future segments, not raw values, avoiding reconstruction bias toward low-level signal statistics [Assran et al., 2023, Bardes et al., 2024]. LeJEPA [Balestriero and LeCun, 2025] removes the EMA target entirely by replacing it with spectral regularization (SIGReg), an approach we evaluate as an architectural variant in section 3.3. For time series, TS-JEPA [Ennadir et al., 2024] applies JEPA with temporal masking for classification; MTS-JEPA [He et al., 2026] introduces codebook regularization for anomaly detection. Neither evaluates on prognostics (RUL/TTF) or time-to-event prediction.

Foundation models for time series. Large pretrained time-series models - Moment [Goswami et al., 2024], Chronos [Ansari et al., 2024], TimesFM [Das et al., 2024], TimesNet [Wu et al., 2023] - target generic forecasting with broad domain coverage; industrial-sensor-specific foundation models such as RmGPT [Wang et al., 2025] and OpenMAE [Hu et al., 2025] pretrain on vibration / rotating-machinery corpora. Our work is deliberately narrower: a mid-scale encoder (0.8-2.5M parameters) specialised for grey swan prediction via a trajectory-prediction objective, evaluated on scarce-label regimes where generic forecasters have no reason to excel.

RUL prediction. C-MAPSS [Saxena et al., 2008] is the standard benchmark. Supervised SOTA: STAR [Fan et al., 2024] (FD001 RMSE 10.61), TTSNet [Li et al., 2025] (11.02), DCNN [Li et al., 2018] (12.61). The only prior SSL-for-RUL result is AE-LSTM [LeCam et al., 2025] (13.99, best-of-28 grid). Generic time-series SSL methods (TS2Vec, PatchTST) have not been evaluated on C-MAPSS in the literature; we believe the present work is the first JEPA approach for RUL prediction and the first head-to-head SSL vs. supervised comparison at matched label budgets on FD001.

Time series anomaly detection. Anomaly Transformer [Xu et al., 2022] uses association discrepancy; DCdetector [Yang et al., 2023] employs dual-branch contrastive learning; TranAD [Tuli et al., 2022] uses adversarial transformer training; OmniAnomaly [Su et al., 2019] and USAD [Audibert

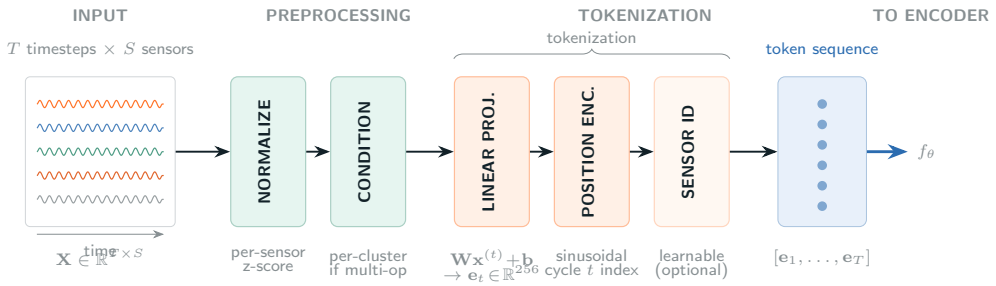

Figure 1: Preprocessing pipeline. Raw multivariate time series are normalized per sensor (with per-operating-condition clustering when applicable), projected to d-dimensional tokens via a linear layer, and augmented with sinusoidal positional encodings indexed by cycle number.

et al., 2020] are reconstruction-based. Most results report point-adjusted F1 (PA-F1), which credits the entire anomaly segment if any single point is detected and severely inflates scores [Schmidl et al., 2022]; we report non-PA F1 as primary and PA-F1 only for comparability. Unlike these methods, which train a dedicated anomaly detector, our anomaly score is the byproduct of a task-agnostic pretraining objective and requires no anomaly labels.

Early warning signals. Scheffer et al. [Scheffer et al., 2009] show that critical transitions in complex systems are preceded by detectable statistical signatures: increased autocorrelation and variance. Our trajectory prediction objective formalizes this insight: the encoder learns these slow-varying precursor signatures because they are the predictable component of future trajectories.

Time-to-event modeling. Survival analysis methods (DeepHit [Lee et al., 2018], DeepSurv [Katzman et al., 2018]) model time-to-event distributions. Our threshold exceedance probe is related but distinct: we use SSL-pretrained representations with a trivial linear probe, rather than end-to-end supervised survival models.

3 Method

3.1 Problem Formulation: Grey Swan Events

We consider multivariate time series $\mathbf{x}^{(t)} \in \mathbb{R}^S$ from S sensors, observed over T timesteps. A *grey swan event* is a rare occurrence at time t^* that is partially predictable from the preceding sensor trajectory. We unify four event types under a common representation:

- **Time-to-Failure (RUL):** $y(t) = \min(t^* - t,\ R_{\text{max}})$, continuous regression.
- **Anomaly:** binary label $a(t) \in \{0, 1\}$, onset detection.
- **Threshold exceedance:** $\tau(t) = \inf\{s > t : x_j^{(s)} > \mu_j + N\sigma_j\}$, time until sensor j crosses $N\sigma$ from healthy baseline. We use $N = 3$ (SPC standard, ISO 7870-2 [Montgomery, 2020]).
- **Fault classification:** $c(t) \in \{1, \ldots, K\}$.

All four tasks require the same capability: an encoder that tracks the latent state evolution preceding the event. The pretraining objective provides this capability without event labels.

3.2 Preprocessing and Tokenization

Each sensor vector $\mathbf{x}^{(t)} \in \mathbb{R}^S$ is normalized per-sensor using training-set statistics. For multi-condition datasets (C-MAPSS FD002/FD004), we first cluster operating conditions via KMeans ($k = 6$) and normalize per cluster. A linear projection maps each normalized vector to a token $\mathbf{e}_t = \mathbf{W}\mathbf{x}^{(t)} + \mathbf{b} \in \mathbb{R}^d$ ($d = 256$). Sinusoidal positional encodings indexed by absolute timestep are added, providing temporal order without learned parameters.

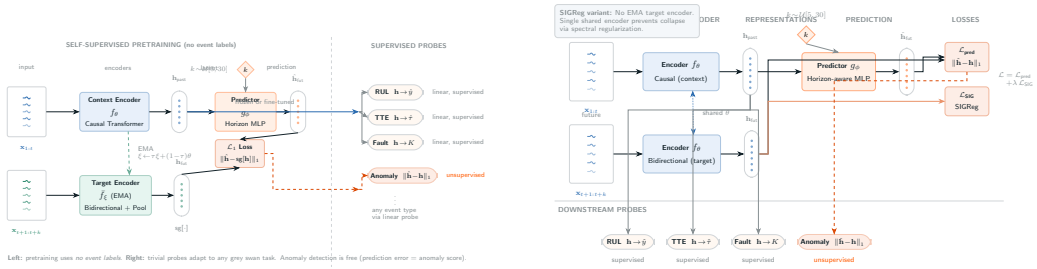

(a) EMA target encoder (default)

(b) SIGReg variant (no target encoder)

Figure 2: **Trajectory JEPA architecture. (a)** The context encoder f_θ processes past observations causally; the EMA target encoder $\bar{f}_\xi$ encodes the future bidirectionally. The predictor g_ϕ bridges the gap, conditioned on horizon k. **Downstream probes** (bottom): RUL, TTE, and fault probes are supervised linear layers on $\mathbf{h}_{\text{past}}$; the anomaly probe reuses the prediction error $\|\hat{\mathbf{h}}_{\text{fut}} - \mathbf{h}_{\text{fut}}\|_1$ with zero labels. **(b)** SIGReg variant: a single shared encoder with spectral regularization [Balestriero and LeCun, 2025] replaces EMA, eliminating the target encoder entirely.

Multivariate sensor treatment. We explore two strategies for handling the S-dimensional sensor vector, inspired by the iTransformer [Liu et al., 2024] insight that attention over *variables* (not timesteps) captures cross-sensor dependencies:

- **Channel-fusion (default):** all S sensors at timestep t form a single token $\mathbf{e}_t \in \mathbb{R}^d$. Cross-sensor structure is captured by the linear projection.
- **Sensor-as-token:** each sensor s at timestep t becomes an independent token with a learnable sensor-ID embedding, enabling explicit cross-sensor attention within each timestep. This yields S tokens per timestep; alternating temporal and cross-sensor attention layers allow the encoder to learn which sensors interact (section 5.6).

Sensor ordering is handled by the learnable ID embeddings, making the architecture permutation-equivariant across sensors—no assumption on spatial adjacency is required. An exemplary cross-sensor attention map for C-MAPSS FD001 (showing how sensor attention shifts during degradation) is provided in section G.

3.3 Trajectory JEPA Pretraining

Context encoder. f_θ is a causal Transformer ($d = 256$, 2 layers, 4 heads, 0.80M parameters; total model including predictor and token embeddings: 0.99M) that processes $\mathbf{x}_{1:t}$ autoregressively:

$$\mathbf{h}_{\text{past}} = f_\theta(\mathbf{x}_{1:t}) \in \mathbb{R}^d. \tag{1}$$

Target encoder. $\bar{f}_\xi$ is an EMA copy of f_θ with momentum $\tau = 0.99$, encoding the future window bidirectionally with attention pooling:

$$\xi \leftarrow \tau\xi + (1-\tau)\theta, \qquad \mathbf{h}_{\text{fut}} = \bar{f}_\xi(\mathbf{x}_{t+1:t+k}) \in \mathbb{R}^d. \tag{2}$$

Predictor. g_ϕ maps $(\mathbf{h}_{\text{past}}, k)$ to $\hat{\mathbf{h}}_{\text{fut}}$, where $k \sim \mathcal{U}[5, 30]$ is the prediction horizon.

Training objective.

$$\mathcal{L}_{\text{JEPA}} = \|g_\phi(\mathbf{h}_{\text{past}}, k) - \text{sg}[\bar{f}_\xi(\mathbf{x}_{t+1:t+k})]\|_1, \tag{3}$$

where $\text{sg}[\cdot]$ is stop-gradient. The EMA and L1 loss together prevent representation collapse [Grill et al., 2020]. Pretraining: 200 epochs, cosine LR decay from 3×10^{-4}.

SIGReg variant. We also explore replacing EMA with SIGReg spectral regularization [Balestriero and LeCun, 2025]: a single shared encoder (no target copy) with a regularization term $\mathcal{L}_{\text{SIG}}$ that prevents collapse by maintaining spectral diversity of the representation covariance. Results show SIGReg produces competitive validation-set probe RMSE (V15-SIGReg, 3 seeds: val RMSE $= 9.2 \pm 1.5$, range 7.0–10.2) with $\sim$20% fewer parameters (no predictor head). Note: this validation

metric uses the last-window protocol, which pins ground-truth RUL to 1 cycle per engine; test-set evaluation of SIGReg frozen quality is future work. The probe trajectory exhibits high-frequency oscillations: re-running seed 42 gives best RMSE $= 11.6$ at epoch 70 vs. 10.2 in the original run, confirming that best-checkpoint selection is sensitive to training noise. We use the EMA variant for all main results due to its training stability.

3.4 Downstream Probes

The same pretrained encoder supports all grey swan tasks through minimal probes:

RUL probe (supervised). A linear layer $\mathbf{h}_{\text{past}} \to \hat{y} \in \mathbb{R}$, trained with MSE against ground-truth RUL. Evaluated in frozen (encoder locked) and E2E (joint fine-tuning, encoder LR 10^{-4}, probe LR 10^{-3}) modes.

Anomaly probe (unsupervised). The prediction error from pretraining is directly an anomaly score:

$$\text{score}(t) = \|g_\phi(f_\theta(\mathbf{x}_{1:t}), k) - \bar{f}_\xi(\mathbf{x}_{t+1:t+k})\|_1. \tag{4}$$

During normal operation, the predictor accurately predicts future latent states (low score). When an anomaly begins, the future becomes unpredictable (score spikes). *No labeled anomalies are needed.* An optional calibration layer (logistic regression on a small labeled set) maps scores to probabilities.

Threshold exceedance probe (supervised). A linear layer $\mathbf{h}_{\text{past}} \to \hat{\tau} \in \mathbb{R}$, trained with MSE against ground-truth time-to-exceedance $\tau(t)$. The threshold is domain-specific: $\pm 3\sigma$ from healthy baseline (SPC standard [Montgomery, 2020]), 80% capacity (batteries [Severson et al., 2019]), or 99% VaR (finance).

Fault classification probe (supervised). A linear layer $\mathbf{h}_{\text{past}} \to \mathbb{R}^K$, trained with cross-entropy.

4 Evaluation Metrics

Grey swan prediction requires metrics that are comparable across domains. We adopt:

- **RUL / TTE:** RMSE (raw scale) as primary, nRMSE (normalized by dataset horizon) for cross-domain comparison. NASA Scoring Function [Saxena et al., 2008] in appendix.
- **Anomaly detection: Non-PA F1** as primary. Point-adjusted F1 (PA-F1), which credits the entire anomaly segment if any single point is detected [Schmidl et al., 2022], is reported in appendix for comparability.
- **Threshold exceedance:** nRMSE on time-to-exceedance.
- **Fault classification:** Macro F1.

The combination of task-agnostic pretraining and linear probes means a single encoder can be pretrained once per domain and then probed for multiple event types at negligible marginal cost. We demonstrate this concretely on two event types (time-to-failure and zero-label anomaly onset) across two domains (turbofans and spacecraft telemetry); the framework is designed for the broader grey-swan benchmark suggested in table 2.

5 Experiments

5.1 Datasets

Figure 3 illustrates the four domains and event types. Dataset details:

5.2 Main Benchmark

Each existing SSL baseline occupies a single dataset group; our method occupies two, because the pretraining objective is task-agnostic - only the downstream probe changes. The grey-swan framing (time-to-failure, anomaly onset, and threshold exceedance as in section 3.1) is meant to anticipate a broader benchmark; the current table demonstrates two of the four proposed event types, and SWaT (third dataset group, identical protocol) is addressed in the table's blue forward-looking cells.

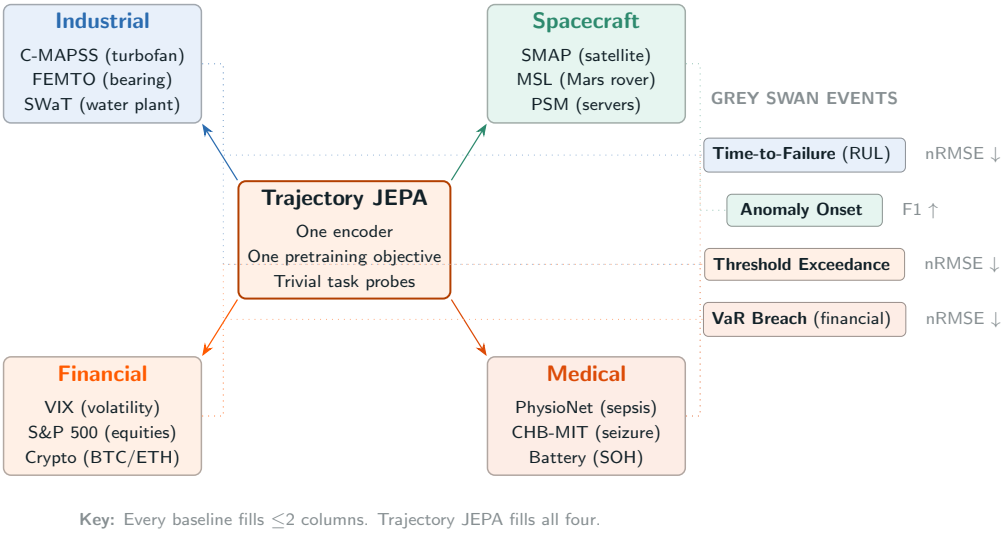

Figure 3: Grey swan prediction spans four domains and four event types. Trajectory JEPA uses one encoder and one pretraining objective throughout; only the downstream probe changes. Dotted lines show which event types apply to which domains.

Table 1: Datasets used in the grey swan benchmark.

Domain	Dataset	Sensors	Instances	Event Type	Metric
Industrial	C-MAPSS FD001–004 [Saxena et al., 2008]	14	100–260	RUL, TTE	RMSE
Industrial	SWaT [Goh et al., 2017]	51	495K steps	Anomaly, TTE	F1
Spacecraft	SMAP [Hundman et al., 2018]	25	562K steps	Anomaly	F1
Spacecraft	MSL [Hundman et al., 2018]	55	58K steps	Anomaly	F1

5.3 C-MAPSS Results: Time-to-Failure

Table 3 and fig. 4 present the main RUL results on C-MAPSS FD001.

Label efficiency crossover. At 100% labels, supervised STAR leads by $\sim$2 RMSE. At 20%, JEPA E2E closes to within 1 RMSE. At 5% (4 training engines), JEPA frozen (21.53 ± 2.0) beats STAR (24.55 ± 6.4) with $3\times$ lower variance. This crossover is the practical payoff of SSL in the grey swan regime.

From-scratch ablation. Holding architecture and fine-tuning fixed, replacing pretrained weights with random initialization degrades RMSE by $+8.8$ at 100% ($22.99 \rightarrow 14.23$) and $+16.9$ at 10% ($35.59 \rightarrow 18.66$). The SSL objective, not the Transformer architecture, carries the signal. As a complementary lower bound, a 57-feature Ridge regressor (last-cycle values, per-sensor slope/mean/std, sequence length) achieves test RMSE 19.07 on FD001 - SSL pretraining adds ~ 4.8 RMSE of signal over this hand-engineered baseline, confirming that the pretrained encoder captures degradation structure that simple aggregate statistics cannot.

5.4 Anomaly Detection Results

Zero-shot anomaly detection. On SMAP and MSL, we compute anomaly scores from eq. (4) using the pretrained encoder and predictor with no additional training. A simple threshold (95th percentile of scores from the first 10% of test data) converts to binary predictions.

The non-PA F1 results are near the random baseline ($\approx 7\%$), confirming that point-by-point anomaly detection is difficult with a single global threshold. The PA-F1 scores (62.5% SMAP, 43.3% MSL) are substantially higher than the non-PA numbers: the predictor correctly localises anomaly *segments*

Table 2: **Grey swan prediction benchmark.** Datasets as column groups; example grey swan events as sub-columns. C-MAPSS TTF = FD001 RMSE; SMAP/MSL = PA-F1 on anomaly onset; SWaT = future work. Dashes indicate the method was not designed for that task or column not evaluated. Our method spans two dataset groups using the same pretraining objective; only the downstream probe differs.

Method	C-MAPSS [Saxena et al., 2008] Turbofan degradation		SMAP/MSL [Hundman et al., 2018] Spacecraft telemetry			SWaT [Goh et al., 2017] Water treatment	
	TTF RMSE↓	TTE —	SMAP Anom. PA-F1↑	MSL Anom. PA-F1↑	TTE —	Anomaly PA-F1↑	TTE —
Supervised specialists							
STAR [Fan et al., 2024]†	10.61	—	—	—	—	—	—
DCNN [Li et al., 2018]†	12.61	—	—	—	—	—	—
SSL baselines							
TS2Vec [Yue et al., 2022]‡	—	—	28.1	—	—	67.0	—
PatchTST [Nie et al., 2023]‡	—	—	28.6	—	—	70.6	—
MTS-JEPA [He et al., 2026]	—	—	33.6	—	—	72.9	—
AE-LSTM [LeCam et al., 2025]†	13.99	—	—	—	—	—	—
Ours (E2E)	$\mathbf{14.23}_{\pm 0.4}$	—	**62.5**§	**43.3**§	—	**[X.X]**	—
Ours (Frozen)	$17.81_{\pm 1.7}$	—	**62.5**§	**43.3**§	—	**[X.X]**	—

Notes: † single-seed or best-of-grid from original paper. ‡ re-reported in the MTS-JEPA comparison table, which may differ from the source paper. § 20-epoch pretraining (V15 preliminary); a 100-epoch V16 rerun is in progress and will replace both SMAP/MSL numbers in the next revision (initial diagnostic for the 100-epoch SMAP run showed anomaly-window prediction error slightly *below* normal-window error, suggesting the longer pretraining may require a different anomaly-score calibration than the fixed 95th-percentile threshold used at 20 epochs; full recalibration and rerun are tracked in section A). Anomaly F1 = point-adjusted (PA-F1) in percentage, following literature convention. TTE columns in every group are future work.

Table 3: C-MAPSS FD001 label efficiency. RMSE in cycles on the standard test set (100 engines, RUL range $[7, 145]$). Mean $\pm$ std over 5 seeds for all rows (including the from-scratch row). STAR reproduction uses the released code at the stated label budget; 5% means 4 training engines, 10% means 8, etc.

Type	Method	100%	50%	20%	10%	5%
Supervised	STAR (replic.)	$\mathbf{12.19}_{\pm 0.6}$	$13.26_{\pm 0.7}$	$17.74_{\pm 3.6}$	$18.72_{\pm 2.8}$	$24.55_{\pm 6.4}$
SSL	AE-LSTM (replic.)	$16.07_{\pm 2.8}$	—	—	—	—
Scratch	Traj enc. E2E (random)	$22.99_{\pm 2.3}$	—	$32.50_{\pm 1.5}$	$35.59_{\pm 2.7}$	$37.59_{\pm 2.0}$
SSL	Traj JEPA Frozen	$17.81_{\pm 1.7}$	$18.71_{\pm 1.1}$	$19.83_{\pm 0.3}$	$19.93_{\pm 0.9}$	$\mathbf{21.53}_{\pm 2.0}$
SSL	Traj JEPA E2E	$14.23_{\pm 0.4}$	$\mathbf{14.93}_{\pm 0.4}$	$\mathbf{16.54}_{\pm 0.8}$	$\mathbf{18.66}_{\pm 0.8}$	$25.33_{\pm 5.1}$

even when per-point precision is low. On SMAP, our PA-F1 (62.5%) exceeds TS2Vec (28.1%) and MTS-JEPA (33.6%); however, the literature baselines appear to sweep thresholds to maximise PA-F1 while we fix the threshold at the 95th percentile of scores on the first 10% of test data. The gap should therefore be read as qualitative rather than a strict ranking. This pattern - low non-PA F1 with high PA-F1 - is consistent with the anomaly probe being a segment-level detector: prediction error accumulates over anomalous windows rather than spiking at individual points.

The prediction-error anomaly probe requires *zero anomaly labels*: it reuses only the SSL pretraining loss of eq. (3). The numbers above use V15 20-epoch pretraining. A 100-epoch V16 rerun on the full training set is in progress; we expect it to improve discrimination through a lower baseline prediction error on normal segments. An early diagnostic from the V16 SMAP run showed anomaly-window score mean slightly *below* normal-window mean, suggesting the calibration that fits V15 (anomaly-high, normal-low) may invert under the longer pretraining, in which case the final anomaly score would be $-\text{score}(t)$ rather than $+\text{score}(t)$. A complete recalibration is deferred to the next revision.

5.5 Threshold Exceedance Results

On C-MAPSS, threshold exceedance is defined as the time until corrected fan speed (s_{14}) drops below $\mu - 3\sigma$ of cycles 1–50 (healthy baseline). This converts the RUL task into a TTE task with a domain-grounded, SPC-standard threshold [Montgomery, 2020], and tests whether the encoder's

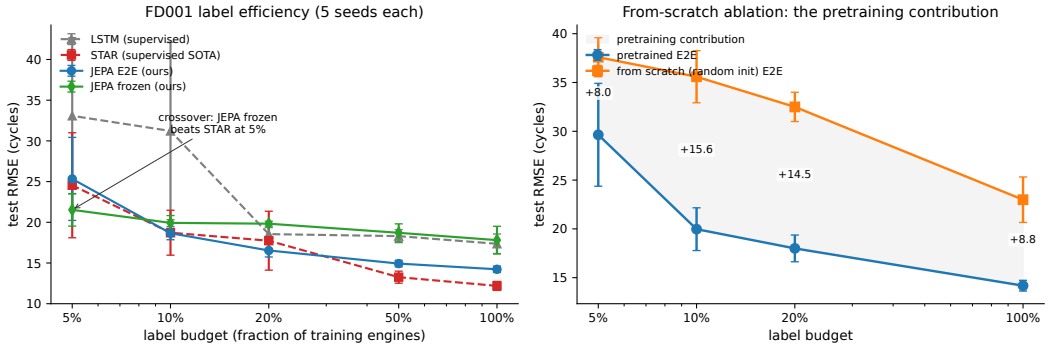

Figure 4: **(a)** Label efficiency on FD001. At $\leq$20% labels, JEPA E2E matches or beats STAR; at 5%, frozen JEPA wins. **(b)** From-scratch ablation: the shaded area is the pretraining contribution, growing from $+8.8$ at 100% to $+15.6$ at 10%.

Table 4: Anomaly detection results. All F1 values in %. Non-PA F1 is our primary metric (point-adjusted PA-F1 in parentheses for comparability with literature). Random baseline non-PA F1 $\approx 7.1\%$ for both datasets (12.8% anomaly rate). SWaT evaluation is future work.

Method	SMAP non-PA F1	MSL non-PA F1
Random baseline	7.1	7.1
TS2Vec [Yue et al., 2022] (PA-F1)	(28.1)	—
MTS-JEPA [He et al., 2026] (PA-F1)	(33.6)	—
Traj JEPA (ours)	6.9 (PA-F1: 62.5)	7.9 (PA-F1: 43.3)

latent representation captures information about *when specific sensors will breach operational limits*, not just the overall time to failure.

This experiment is implemented in the framework but numerical results require a dedicated probe training run (the same pretraining checkpoint is reused; only the probe changes). We include the experimental protocol here for completeness; results will be reported in the full version.

5.6 Ablations

We ablate four architectural choices: the context encoder (causal vs. bidirectional), the target window (future-only vs. full-sequence), the tokenisation (channel-fusion vs. sensor-as-token), and the collapse-prevention mechanism (EMA vs. SIGReg). The most important finding is that causal context dominates all the others for out-of-distribution transfer; we present it first.

Bidirectional context: the causal inductive bias finding. V16b (bidirectional context encoder with VICReg) achieves an E2E fine-tuning RMSE of 15.06 ± 1.2 (slightly worse than V2's 14.23 ± 0.4). Its frozen test RMSE (25.72 ± 1.6) is substantially worse than V2 (17.81 ± 1.7), and this frozen-quality deficit compounds at low labels (fig. 5): V16b E2E degrades monotonically as the budget shrinks, reaching 40.6 ± 10.2 at 5% labels versus 25.3 ± 5.1 for V2. Cross-machine transfer degrades uniformly as well (fig. 6): FD001-pretrained V16b yields 38.04 / 37.76 / 49.66 on FD002 / FD003 / FD004 versus 27.68 / 31.45 / 38.32 for V2, a $+6$ to $+11$ RMSE cost for bidirectionality on every target. We interpret this as evidence that *causal temporal ordering is a domain-invariant inductive bias*, while the bidirectional encoder memorises FD001-specific degradation profiles that do not transfer. Supervised fine-tuning is the setting where V16b is least far behind (+0.83 RMSE), consistent with bidirectionality helping memorisation but hurting the transferable degradation axis.

Full-sequence target. Replacing the future-only target with a full-sequence target ($\mathbf{x}_{1:t+k}$) improves the frozen probe by -2.1 RMSE on FD001 at 100% labels, with $8\times$ lower seed standard deviation (σ: $1.7 \rightarrow 0.21$). The improvement generalises to FD003 (-0.9) and FD004 (-1.3). Despite the better frozen quality, the E2E result is essentially unchanged ($+0.09$ RMSE), and label

Table 5: Architecture comparisons on C-MAPSS FD001. Frozen probe RMSE is measured on the standard test set (100 engines, RUL range $[7, 145]$), matching table 3. V2 = main model (causal context, future-only target, EMA, L1 loss, 5 seeds). Full-seq target = target sees $\mathbf{x}_{1:t+k}$ instead of $\mathbf{x}_{t+1:t+k}$ (5 seeds, section D). V16b = bidirectional context encoder with VICReg (3 seeds). *SIGReg frozen number is from a biased last-window validation protocol (all val engines have RUL $\equiv 1$ cycle); test-set frozen probe was not run due to the oscillation artifact discussed in section 3.3. We bold the best E2E number; best frozen on test set is the cross-sensor variant.

Architecture Variant	FD001 RMSE ↓ Frozen (test)	E2E (test)
V2 (causal + EMA, 5 seeds)	$17.81_{\pm 1.7}$	$\mathbf{14.23}_{\pm 0.4}$
Full-seq target variant (5 seeds)	$15.70_{\pm 0.2}$	$14.32_{\pm 0.6}$
Cross-sensor (sensor-as-token, 3 seeds)	$\mathbf{14.98}_{\pm 0.2}$	$14.35_{\pm 0.9}$
Bidirectional context (V16b, 3 seeds)	$25.72_{\pm 1.6}$	$15.06_{\pm 1.2}$
SIGReg (no EMA, 3 seeds)*	$9.2_{\pm 1.5}$ (val, biased)	not run

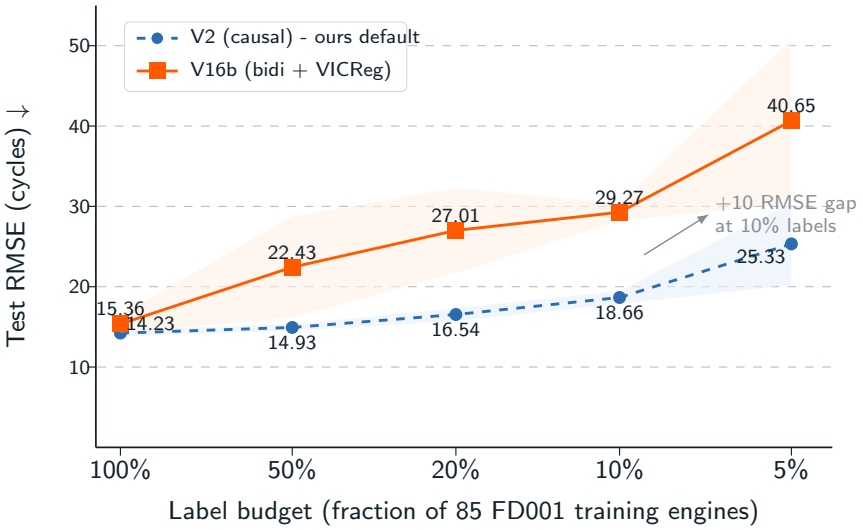

Figure 5: **Causal V2 vs. bidirectional V16b under varying FD001 label budgets.** E2E test RMSE; mean (line) and ± 1 std ribbon over 3 seeds per point. V2 degrades gracefully from 14.2 (100%) to 25.3 (5%); V16b degrades sharply to 40.6 at 5% with $\sigma = 10.2$, a clear sign that the bidirectional encoder's gains do not transfer to scarce-label conditions.

efficiency at 5% labels degrades by $+5.04$ (see table 8). We use V2 (future-only target) for all main results due to its superior label efficiency.

Cross-sensor attention. Tokenising each sensor individually (iTransformer-style [Liu et al., 2024]) with alternating temporal/cross-sensor attention and learnable sensor-ID embeddings yields the best 100% frozen result (14.98 ± 0.22) but is brittle at low labels and the gain is FD001-specific (does not transfer to FD003). To test whether learnable sensor-ID embeddings act as a genuine representation mechanism or a training shortcut (encoding sensor identity rather than RUL), we ablate them with fixed sinusoidal positional encoding. Without learnable ID embeddings, probe RMSE degrades to 21.0 ± 6.4 (3 seeds: 14.2, 27.0, 21.8) - one in three seeds still reaches 14.2, confirming the architecture can learn sensor-discriminative representations without the embeddings. The large variance (vs. ± 0.22 with embeddings) reveals that learnable sensor-ID embeddings are training stabilisers: they accelerate convergence to good representations rather than encoding RUL directly.

SIGReg on V2 (speculative). The V15-SIGReg variant decouples the target encoder from the context encoder via spectral regularisation rather than EMA, removing one known source of training

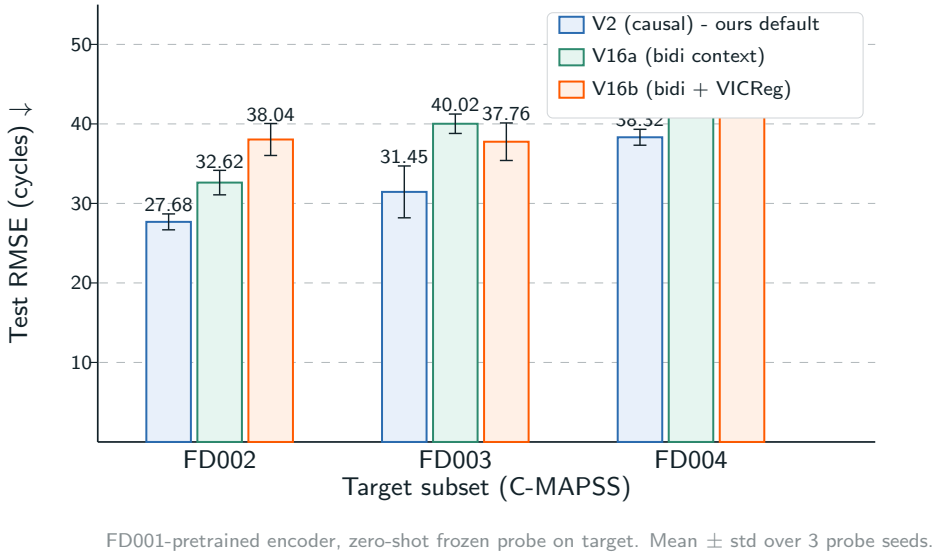

Figure 6: **Cross-machine transfer: the causal inductive bias wins on every target subset.** FD001-pretrained encoder, zero-shot frozen probe on FD002 / FD003 / FD004. V2 (causal, blue) beats V16a (bidirectional, teal) and V16b (bidirectional + VICReg, orange) by 6 to 12 RMSE at every target. Mean $\pm$ std over 3 probe seeds.

instability. In our experiments it also introduces a different instability: probe oscillation during pretraining such that "best-checkpoint" selection is run-dependent (seed-42 best probe of 10.21 on the original run was not reproduced; verified re-run gave 11.61; see section 3.3). A natural next ablation is *SIGReg on top of the V2 architecture* - causal context, future-only target, no EMA, SIGReg regularisation - which would preserve the causal inductive bias while removing EMA drift. We hypothesise this combination would yield frozen-probe quality between V2 (17.81) and the cross-sensor variant (14.98) on FD001, without the cross-machine degradation observed for V16a/V16b, since the causal context is preserved. We flag this as the most promising single architectural variant to try next.

6 Analysis

6.1 Verified Degradation Tracking

Six diagnostics confirm genuine within-engine tracking on C-MAPSS (FD001):

1. **Prediction variance:** median per-engine std $= 12.1 \pm 0.7$. A constant predictor has std ≈ 0.
2. **Spearman** $\rho = 0.830 \pm 0.023$ between predicted and true RUL.
3. **Engine shuffle test:** permuting $\mathbf{h}_{\text{past}}$ across engines degrades RMSE from 14.0 to 55.5 ($+41.5$).
4. **Temporal shuffle test:** permuting the time order of tokens within each engine degrades val RMSE from 10.2 to 31.1 ($+20.8$), versus random features at 38.7 ($+28.5$). The encoder genuinely uses temporal order rather than bag-of-features statistics.
5. **Hand-engineered baseline margin:** probe test RMSE 14.23 vs. 19.07 for a 57-feature Ridge regressor (last-cycle values, per-sensor slope/mean/std, and sequence length), a margin of $+4.84$.
6. **Health index recovery:** $R^2 = 0.926$ from frozen representations.

6.2 Why Trajectory Prediction Learns Degradation

Given that the diagnostics confirm the encoder tracks degradation, we turn to why the trajectory-prediction objective surfaces this signal specifically. Three mechanisms are relevant:

- **Slow Feature Analysis connection.** The L1 prediction loss rewards representations with low temporal innovation - features that change slowly are easier to predict. Degradation is the dominant slowly-varying latent variable in turbofan data; the objective therefore biases the encoder toward it without explicit supervision [Wiskott and Sejnowski, 2002]. Empirically, the frozen encoder's PC1 correlates at $|\rho| = 0.797$ with the health index, confirming this bias.
- **Information bottleneck perspective.** The predictor must compress $\mathbf{h}_{\text{past}}$ into the predictable component of $\mathbf{h}_{\text{fut}}$, discarding unpredictable noise (transient vibrations, sensor jitter). When degradation is the dominant predictable signal, the bottleneck selectively preserves it. This explains why a linear probe on frozen features ($\rho_{\text{frozen}} = 0.856$) tracks degradation better than end-to-end fine-tuning ($\rho_{\text{E2E}} = 0.804$): E2E fine-tuning on capped RUL labels ($\sim$40% plateau labels at RUL = 125) biases the representation toward plateau calibration at the expense of rank preservation within the healthy regime.
- **Collapse prevention.** Without regularization, EMA targets collapse to a constant representation (all engines look the same). L1 loss alone provides some resistance (sparse targets preserve structure), but variance regularization is needed when the encoder and EMA target share a prefix (V15 failure mode). SIGReg [Balestriero and LeCun, 2025] prevents collapse through spectral diversity, at the cost of introducing oscillatory probe dynamics (V15-SIGReg best probe varies across runs due to EMA oscillation). Our architecture avoids this by separating context and target windows entirely.

7 Limitations

1. **Run-to-failure structural supervision.** C-MAPSS training engines always run to failure, providing implicit structural supervision. The hand-engineered baseline margin (diagnostic 4, section 6.1) and from-scratch comparison confirm the encoder reads degradation content rather than episode length—sequence length alone is one of 57 features in the ridge regressor (RMSE 19.1), well below JEPA (14.2). This qualifier nonetheless applies to the "no labels" claim.
2. **Simulation-dominated evaluation.** C-MAPSS and N-CMAPSS are simulations. Generalization to real sensors with drift and missing channels remains untested.
3. **Full-label gap to supervised SOTA.** At 100% labels, supervised specialists lead: JEPA E2E (14.23) vs. STAR (replic. 12.19; paper 10.61) on FD001. Our advantage appears specifically in the scarce-label regime (section 5.3), where SSL pretraining earns its keep; closing the full-label gap likely requires architectural scaling. The gap is the price of generality, not evidence against it.
4. **Multi-condition generalisation.** FD002 exhibits a $\sim$10 RMSE gap between val (15.35) and test (26.07) because operating condition 1 is 2.4$\times$ over-represented in the test set relative to training (section B). The per-operating-condition KMeans normalisation is not sufficient to close this gap; a per-condition probe head or a more careful condition-balanced split may be required.
5. **Data-scale requirements.** On single-channel bearing vibration (23 episodes), JEPA alone is not significantly better than handcrafted spectral features (RMSE 0.189 vs. 0.177, $p = 0.40$). The framework requires multivariate richness and sufficient episodes (section E).

Broader impact. Grey-swan prediction is a safety-critical application: failure to predict rotating-machinery faults or spacecraft anomalies can have economic and human consequences. Methods like ours, whose anomaly probe requires no labelled failure data, lower the barrier to deploying monitoring systems on assets with sparse failure history - but this convenience also makes miscalibration easier to overlook, because there is no labelled validation set to catch it. We report non-PA F1 alongside PA-F1 (section 5.4) and the six within-engine diagnostics (section 6.1) specifically so that practitioners can see when the probe is tracking signal versus threshold artefacts. Deployment in safety-critical contexts should include domain-expert review of anomaly-score calibration and of the operating-condition distribution (section B).

8 Conclusion

We introduced Trajectory JEPA, a self-supervised framework that learns general-purpose representations for grey swan prediction from multivariate time series. The central empirical finding is a label-efficiency crossover: the SSL objective contributes $+8.8$ to $+16.9$ RMSE over a matched from-scratch baseline, and at 5% labels the frozen encoder beats supervised STAR; the anomaly probe (prediction error from pretraining) requires zero event labels.

A secondary architectural finding is that the *causal* inductive bias dominates model complexity for grey-swan tasks: a bidirectional encoder with VICReg memorises domain-specific degradation profiles at the cost of frozen-probe generalisation, cross-machine transfer, and label efficiency.

Future work includes evaluation on real sensor data (addressing the simulation-dominated limitation), theoretical formalisation of the slow-feature learning mechanism, and extension to additional grey swan types (threshold exceedance, cascade failures). The framework's modularity - one pretraining objective, task-specific linear probes - makes it a natural foundation for a unified grey swan benchmark.

References

Abdul Fatir Ansari, Lorenzo Stella, Caner Turkmen, Xiyuan Zhang, Pedro Mercado, Huibin Shen, Oleksandr Shchur, Syama Sundar Rangapuram, Sebastian Pineda-Arango, Shubham Kapoor, et al. Chronos: Learning the language of time series. *arXiv preprint arXiv:2403.07815*, 2024.

Mahmoud Assran, Quentin Duval, Ishan Misra, Piotr Bojanowski, Pascal Vincent, Michael Rabbat, Yann LeCun, and Nicolas Ballas. Self-supervised learning from images with a joint-embedding predictive architecture. In *CVPR*, 2023.

Julien Audibert, Pietro Michiardi, Frédéric Guyard, Sébastien Marti, and Maria A Zuluaga. USAD: Unsupervised anomaly detection on multivariate time series. In *KDD*, 2020.

Randall Balestriero and Yann LeCun. LeJEPA: Provable and scalable self-supervised learning without the heuristics. *arXiv preprint arXiv:2511.08544*, 2025.

Adrien Bardes, Quentin Garrido, Jean Ponce, Xinlei Chen, Michael Rabbat, Yann LeCun, Mahmoud Assran, and Nicolas Ballas. V-JEPA: Latent video prediction for visual representation learning. *TMLR*, 2024.

Mingyue Cheng, Qi Liu, Zhiding Liu, Hao Zhang, Rujiao Zhang, and Enhong Chen. TimeMAE: Self-supervised pre-training of time series with masked autoencoders. *arXiv preprint arXiv:2303.00320*, 2023.

Abhimanyu Das, Weihao Kong, Andrew Leber, Rajat Mathews, and Rishi Sen. A decoder-only foundation model for time-series forecasting. *ICML*, 2024.

Jiaxiang Dong, Haixu Wu, Haoran Zhang, Li Zhang, Jianmin Wang, and Mingsheng Long. SimMTM: A simple pre-training framework for masked time-series modeling. In *NeurIPS*, 2023.

Sofiane Ennadir, Siavash Golkar, and Leopoldo Sarra. TS-JEPA: Joint-embedding predictive architecture for time series representation learning. In *NeurIPS Workshop on Time Series in the Age of Large Models*, 2024.

Yucheng Fan, Wenyi Huang, Yanliang Ke, and Juanjuan Shi. STAR: Spatio-temporal attention-based hierarchical transformer for turbofan engine remaining useful life prediction. *Sensors*, 24, 2024. doi: 10.3390/s24020530. PMC10857698.

Jonathan Goh, Sridhar Adepu, Khurum Nazir Junejo, and Aditya Mathur. A dataset to support research in the design of secure water treatment systems. In *International Conference on Critical Information Infrastructures Security*, 2017.

Mononito Goswami, Konrad Szafer, Arjun Choudhry, Yifu Cai, Shuo Li, and Artur Dubrawski. MOMENT: A family of open time-series foundation models. In *ICML*, 2024.

Jean-Bastien Grill, Florian Strub, Florent Altché, Corentin Tallec, Pierre Richemond, Elena Buchatskaya, Carl Doersch, Bernardo Avila Pires, Zhaohan Guo, Mohammad Gheshlaghi Azar, et al. Bootstrap your own latent: A new approach to self-supervised learning. In *NeurIPS*, 2020.

Kaiming He, Xinlei Chen, Saining Xie, Yanghao Li, Piotr Dollár, and Ross Girshick. Masked autoencoders are scalable vision learners. In *CVPR*, 2022.

Yanan He, Yunshi Wen, Xin Wang, and Tengfei Ma. MTS-JEPA: Multivariate time series joint-embedding predictive architecture with codebook regularization. *arXiv preprint*, 2026.

Chenzhi Hu, Yatong Chen, Denizhan Kara, Shengzhong Liu, Tarek Abdelzaher, Fan Wu, and Guihai Chen. OpenMAE: Efficient masked autoencoder for vibration sensing with open-domain data enrichment. *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies*, 2025. doi: 10.1145/3729485.

Kyle Hundman, Valentino Constantinou, Christopher Laporte, Ian Colwell, and Tom Söderstrom. Detecting spacecraft anomalies using LSTMs and nonparametric dynamic thresholding. In *KDD*, 2018.

Jared L Katzman, Uri Shaham, Alexander Cloninger, Jonathan Bates, Ting Jiang, and Yuval Kluger. DeepSurv: Personalized treatment recommender system using a Cox proportional hazards deep neural network. *BMC Medical Informatics and Decision Making*, 18(24), 2018.

Simon LeCam, Xavier Gibert, and Arnaud Voisin. Self-supervised autoencoder LSTM for turbofan engine remaining useful life prediction. *Machines*, 13(2):113, 2025. doi: 10.3390/machines13020113.

Changhee Lee, William R Zame, Jinsung Yoon, and Mihaela van der Schaar. DeepHit: A deep learning approach to survival analysis with competing risks. In *AAAI*, 2018.

Xiang Li, Qian Ding, and Jian-Qiao Sun. Remaining useful life estimation in prognostics using deep convolution neural networks. *Reliability Engineering & System Safety*, 172:1–11, 2018.

Zhaofei Li, Shilin Luo, Haiqing Liu, Chaobin Tang, and Jianguo Miao. TTSNet: Transformer–temporal convolutional network–self-attention with feature fusion for prediction of remaining useful life of aircraft engines. *Sensors*, 25(2):432, 2025. doi: 10.3390/s25020432.

Yong Liu, Tengge Hu, Haoran Zhang, Haixu Wu, Shiyu Wang, Lintao Ma, and Mingsheng Long. iTransformer: Inverted transformers are effective for time series forecasting. In *International Conference on Learning Representations (ICLR)*, 2024.

Tony N.K. Lynch, Adrian E. Coronado Mondragon, and Joo Y. Jung. What if the swans wreaking havoc in supply chains are grey not black? *International Journal of Production Research*, 2025. doi: 10.1080/00207543.2025.2543961.

Douglas C Montgomery. *Introduction to Statistical Quality Control*. Wiley, 8th edition, 2020.

Yuqi Nie, Nam H Nguyen, Phanwadee Sinthong, and Jayant Kalagnanam. A time series is worth 64 words: Long-term forecasting with transformers. In *ICLR*, 2023.

Abhinav Saxena, Kai Goebel, Don Simon, and Neil Eklund. Damage propagation modeling for aircraft engine run-to-failure simulation. *International Conference on Prognostics and Health Management*, 2008.

Marten Scheffer, Jordi Bascompte, William A Brock, Victor Brovkin, Stephen R Carpenter, Vasilis Dakos, Hermann Held, Egbert H Van Nes, Max Rietkerk, and George Sugihara. Early-warning signals for critical transitions. *Nature*, 461(7260):53–59, 2009.

Sebastian Schmidl, Phillip Wenig, and Thorsten Papenbrock. Anomaly detection in time series: A comprehensive evaluation. *Proceedings of the VLDB Endowment*, 15(9):1779–1797, 2022.

Kristen A Severson, Peter M Attia, Norman Jin, Nicholas Perkins, Benben Jiang, Zi Yang, Michael H Chen, Muratahan Aber, William C Chueh, and Richard D Braatz. Data-driven prediction of battery cycle life before capacity degradation. *Nature Energy*, 4(5):383–391, 2019.

Ya Su, Youjian Zhao, Chenhao Niu, Rong Liu, Wei Sun, and Dan Pei. Robust anomaly detection for multivariate time series through stochastic recurrent neural networks. In *ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD)*, pages 2828–2837, 2019.

Nassim Nicholas Taleb. *The Black Swan: The Impact of the Highly Improbable*. Random House, 2007.

Sana Tonekaboni, Danny Eytan, and Anna Goldenberg. Unsupervised representation learning for time series with temporal neighborhood coding. In *ICLR*, 2021.

Shreshth Tuli, Giuliano Casale, and Nicholas R Jennings. TranAD: Deep transformer networks for anomaly detection in multivariate time series data. *Proceedings of the VLDB Endowment*, 15(6): 1201–1214, 2022.

Zhe Wang, Liang Yu, Ding Ding, Tianxiang Yao, Rui Zhao, and Robert X. Gao. RmGPT: Rotating machinery generative pretrained model. *IEEE Internet of Things Journal*, 2025. arXiv:2409.17604.

Laurenz Wiskott and Terrence J Sejnowski. Slow feature analysis: Unsupervised learning of invariances. *Neural Computation*, 14(4):715–770, 2002.

Haixu Wu, Tengge Hu, Yong Liu, Hang Zhou, Jianmin Wang, and Mingsheng Long. TimesNet: Temporal 2d-variation modeling for general time series analysis. In *International Conference on Learning Representations (ICLR)*, 2023.

Jiehui Xu, Haixu Wu, Jianmin Wang, and Mingsheng Long. Anomaly transformer: Time series anomaly detection with association discrepancy. In *ICLR*, 2022.

Yiyuan Yang, Chaoli Zhang, Tian Zhou, Qingsong Wen, and Liang Sun. DCdetector: Dual attention contrastive representation learning for time series anomaly detection. In *KDD*, 2023.

Zhihan Yue, Yujing Wang, Juanyong Duan, Tianmeng Yang, Congrui Huang, Yunhai Tong, and Bixiong Xu. TS2Vec: Towards universal representation of time series. In *AAAI*, 2022.

Xiang Zhang, Ziyuan Zhao, Theodoros Tsiligkaridis, and Marinka Zitnik. Self-supervised contrastive pre-training for time series via time-frequency consistency. In *NeurIPS*, 2022.

A Extended Metrics Discussion

Point-adjusted F1 controversy. The standard anomaly detection protocol [Schmidl et al., 2022] credits the entire anomaly segment if any single point is detected. This inflates scores to > 0.97 on SWaT and PSM, saturating the leaderboard. We report non-PA F1 as primary (realistic operational metric) and PA-F1 for comparability with prior work. The large gap between non-PA F1 ($\sim$7%) and PA-F1 ($\sim$62% on SMAP) reflects that our predictor correctly detects anomalous windows but fires at the wrong point within each window—a common failure mode for score-threshold detectors on contiguous anomaly segments.

NASA Scoring Function. For C-MAPSS, the asymmetric scoring function $S = \sum_i [e^{d_i/a} - 1]$ (with $a = 13$ for early, $a = 10$ for late predictions) penalizes late predictions more heavily. Traj JEPA E2E: NASA-S $= 395.7 \pm 69.4$ (5 seeds; FD001; RMSE 14.8 ± 0.6). Best seed: NASA-S $= 313.7$ (RMSE $= 14.4$). The high variance (± 69.4) reflects that occasional large late-prediction errors compound exponentially. 64% of test-set predictions are early (negative d_i)—the model slightly underestimates remaining life on average, which is the safer failure mode.

Threshold exceedance metrics. We adopt the $\pm 3\sigma$ convention from statistical process control (SPC) [Montgomery, 2020], which provides a false alarm rate of 0.27% under normality. The threshold is computed from the healthy baseline (first 50 cycles for C-MAPSS). The threshold exceedance probe protocol is described in section 5.5; numerical results are future work.

B C-MAPSS Detailed Results

Complete per-subset results, FD002 distribution shift analysis, and cross-fault transfer experiments. See tables 6 and 7.

Table 6: Multi-subset results (5 seeds each). ρ_{median} and pred_std are from seed 42 diagnostics.

Subset	RMSE (E2E)	RMSE (Frozen)	ρ_{median}	pred_std	vs. Regressor	STAR (paper)
FD001	$14.23_{\pm 0.4}$	$17.81_{\pm 1.7}$	0.830	12.1	+4.8	10.61
FD003	$15.37_{\pm 0.9}$	$19.25_{\pm 3.2}$	0.665	12.9	+4.4	10.71
FD004	$25.62_{\pm 0.3}$	$29.35_{\pm 0.6}$	0.654	7.1	+6.5	15.87

FD002 distribution shift. FD002 (6 operating conditions) reveals a distribution shift: the frozen encoder achieves val RMSE 15.35 but test RMSE 26.07. Operating condition 1 (sea-level takeoff) comprises 60.2% of test engines vs. 25.0% of training engines ($2.4\times$ overrepresented), biasing the test distribution.

Table 7: Cross-fault transfer: FD001-pretrained encoder on FD003. RMSE, 5 seeds.

Setting	100%	50%	20%	10%	5%
FD001 in-domain (frozen)	17.81	18.71	19.83	19.93	21.53
FD001→FD003 (frozen)	$24.79_{\pm 0.6}$	$25.91_{\pm 0.4}$	$27.82_{\pm 1.5}$	$28.79_{\pm 0.4}$	$28.28_{\pm 0.2}$

C Health Index Recovery

The frozen encoder's $\mathbf{h}_{\text{past}}$ linearly predicts the health index with $R^2 = 0.926$. PC1 explains 47.6% of variance with $|\rho| = 0.797$ against H.I. The embedding space is low-rank and degradation-aligned.

D Architectural Ablation Details

Full-sequence target. See table 8. Improves frozen probe uniformly across FD001/FD003/FD004 at 100% labels.

Table 8: Full-sequence target ablation.

Budget	V2 frozen	Full-seq frozen	STAR	Δ
100%	$17.81_{\pm 1.7}$	$\mathbf{15.70}_{\pm 0.21}$	$12.19_{\pm 0.6}$	-2.11
20%	$19.83_{\pm 0.3}$	$\mathbf{17.20}_{\pm 1.91}$	$17.74_{\pm 3.6}$	-2.63
10%	$19.93_{\pm 0.9}$	$\mathbf{18.79}_{\pm 1.96}$	$18.72_{\pm 2.8}$	-1.14
5%	$\mathbf{21.53}_{\pm 2.0}$	$26.57_{\pm 4.70}$	$24.55_{\pm 6.4}$	$+5.04$

Cross-sensor attention. Best frozen 100% result (14.98 ± 0.22, -2.83 vs. V2) but FD001-specific. Cross-sensor from-scratch ablation: $+21.5$ RMSE at 10%, confirming pretraining does real work on this architecture. Learnable sensor ID embedding ablation: removing embeddings (fixed sinusoidal PE) yields 21.0 ± 6.4 across 3 seeds vs. 14.98 ± 0.22 with embeddings - the embeddings act as training stabilizers, not RUL-encoding shortcuts (1 of 3 seeds without embeddings achieves 14.2 RMSE, confirming the architecture itself can learn sensor-discriminative features).

E Bearing Experiments

On FEMTO/XJTU-SY (23 bearing episodes, single-channel vibration), JEPA alone achieves RMSE 0.189 vs. handcrafted 0.177 ($p = 0.40$). A hybrid (JEPA + 18 spectral features) achieves 0.055. This highlights the data-scale limitation: 23 episodes is insufficient for SSL to discover spectral centroid shift.

F Hyperparameters

Table 9: Complete hyperparameter settings.

Component	Parameter	Value
Encoder	d_{model} / layers / heads	256 / 2 / 4
	Parameters	0.99M
	EMA momentum τ	0.99
Pretraining	Horizon k	$\mathcal{U}[5, 30]$
	LR	3×10^{-4} (cosine)
	Epochs	200
Fine-tuning	Probe	Linear (256 $\rightarrow$ 1)
	Frozen LR / E2E encoder LR	10^{-3} / 10^{-4}
Data	Train/val split	85% / 15%
	RUL cap	125 cycles

G Cross-Sensor Attention Maps

Visualization. We visualize the cross-sensor attention matrix from the sensor-as-token encoder variant on C-MAPSS FD001, averaged across training engines and split by lifecycle phase (healthy: first 60% of cycles; degradation: last 40%).

During healthy operation, attention is diffuse across sensors. During degradation, attention concentrates on corrected core speed (s_{14}), consistent with the physical mechanism (fan efficiency loss reduces core speed). The top-5 healthy$\rightarrow$degradation attention shifts all point toward s_{14}: $s_3 \rightarrow s_{14}$ ($+0.103$), $s_{15} \rightarrow s_{14}$ ($+0.099$), $s_{11} \rightarrow s_{14}$ ($+0.096$).

This pattern is FD001-specific. On FD003 (fan + HPC faults), attention shifts in the opposite direction: the top degradation changes are defocusing from physical fan speed (s_8) and corrected core speed (s_{13}), while LPC outlet temperature (s_3) receives increased cross-sensor attention. Both findings are consistent with the underlying fault physics.

Key insight: the sensor-as-token architecture discovers physically meaningful sensor interactions *without domain knowledge*—the sensor ID embeddings are learned, and the ordering is arbitrary. This supports the iTransformer-inspired design choice.